

```
In [3]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
```

```
In [4]: beneficiary = pd.read_csv("/content/sample_data/Beneficiarydata.csv")
in_patient = pd.read_csv("/content/sample_data/Inpatientdata.csv")
out_patient = pd.read_csv("/content/sample_data/Outpatientdata.csv")
provider = pd.read_csv("/content/sample_data/Provider.csv")
```

```
In [4]:
```

Checking for null values in all the tables

```
In [5]: tables = {
    "Beneficiary": beneficiary,
    "In_patient": in_patient,
    "Out_patient": out_patient,
    "Provider": provider
}

# Loop through each dataframe and check for null values
for name, df in tables.items():
    print(f"\nNull values in {name} table:")
    print(df.isnull().sum())
```

Null values in Beneficiary table:

BeneID	0
DOB	0
DOD	63394
Gender	0
Race	0
RenalDiseaseIndicator	0
State	0
County	0
NoOfMonths_PartACov	0
NoOfMonths_PartBCov	0
ChronicCond_Alzheimer	0
ChronicCond_Heartfailure	0
ChronicCond_KidneyDisease	0
ChronicCond_Cancer	0
ChronicCond_ObstrPulmonary	0
ChronicCond_Depression	0
ChronicCond_Diabetes	0
ChronicCond_IschemicHeart	0
ChronicCond_Osteoporosis	0
ChronicCond_rheumatoidarthritis	0
ChronicCond_stroke	0
IPAnnualReimbursementAmt	0
IPAnnualDeductibleAmt	0
OPAnnualReimbursementAmt	0
OPAnnualDeductibleAmt	0

dtype: int64

Null values in In_patient table:

BeneID	0
ClaimID	0
ClaimStartDt	0
ClaimEndDt	0
Provider	0
InscClaimAmtReimbursed	0
AttendingPhysician	31
OperatingPhysician	3962
OtherPhysician	8538
AdmissionDt	0
ClmAdmitDiagnosisCode	0
DeductibleAmtPaid	196
DischargeDt	0
DiagnosisGroupCode	0
ClmDiagnosisCode_1	0
ClmDiagnosisCode_2	54
ClmDiagnosisCode_3	169
ClmDiagnosisCode_4	404
ClmDiagnosisCode_5	719
ClmDiagnosisCode_6	1197
ClmDiagnosisCode_7	1736
ClmDiagnosisCode_8	2360
ClmDiagnosisCode_9	3238
ClmDiagnosisCode_10	8664
ClmProcedureCode_1	4118
ClmProcedureCode_2	8297
ClmProcedureCode_3	9328
ClmProcedureCode_4	9522
ClmProcedureCode_5	9549
ClmProcedureCode_6	9551

dtype: int64

Null values in Out_patient table:

BeneID	0
ClaimID	0
ClaimStartDt	0
ClaimEndDt	0
Provider	0
InscClaimAmtReimbursed	0
AttendingPhysician	316
OperatingPhysician	104237
OtherPhysician	78222
ClmDiagnosisCode_1	2578
ClmDiagnosisCode_2	47731
ClmDiagnosisCode_3	76575
ClmDiagnosisCode_4	95371
ClmDiagnosisCode_5	107875
ClmDiagnosisCode_6	114035
ClmDiagnosisCode_7	117871
ClmDiagnosisCode_8	120310
ClmDiagnosisCode_9	122278
ClmDiagnosisCode_10	125578
ClmProcedureCode_1	125807
ClmProcedureCode_2	125832
ClmProcedureCode_3	125839
ClmProcedureCode_4	125841
ClmProcedureCode_5	125841
ClmProcedureCode_6	125841
DeductibleAmtPaid	0
ClmAdmitDiagnosisCode	100036

dtype: int64

Null values in Provider table:

Provider 0
dtype: int64

In [6]: *# Here we can there are no null values in primary key cols of the table*

In [7]: *# These are the required features to be included.*

Target Variable: Flag claims as fraudulent or non-fraudulent (based on provide

Features to include:

- # Claim amount (InscClaimAmtReimbursed)*
- # Provider ID*
- # Number of physicians involved*
- # Chronic conditions*
- # Claim type (inpatient / outpatient)*

In [8]: *# In order to have these features we need a summary table where we can only have*

In [9]: *# Adding a claim type column to both the in_patient and out_patient table*

```
in_patient["ClaimType"] = "Inpatient"
out_patient["ClaimType"] = "Outpatient"
```

In [10]: *# Adding a physician count column to both the in_patient and out_patient table*

```
def count_physicians(row):
    return sum([pd.notna(row.get("AttendingPhysician")),
```

```

        pd.notna(row.get("OperatingPhysician")),
        pd.notna(row.get("OtherPhysician"))
    ])

```

```

in_patient["NumPhysicians"] = in_patient.apply(count_physicians, axis=1)
out_patient["NumPhysicians"] = out_patient.apply(count_physicians, axis=1)

```

```

In [11]: # Selecting the relevant columns
in_patient_sel = in_patient[["BeneID", "ClaimID", "InscClaimAmtReimbursed", "Pro
out_patient_sel = out_patient[["BeneID", "ClaimID", "InscClaimAmtReimbursed", "P

```

```

In [12]: # Combine inpatient and outpatient
claims = pd.concat([in_patient_sel, out_patient_sel], ignore_index=True)

```

```

In [13]: # Here we now have a combined claims table from both in_patient and out_patient
claims.head()

```

```

Out[13]:

```

	BeneID	ClaimID	InscClaimAmtReimbursed	Provider	NumPhysicians	ClaimTy
0	BENE11014	CLM67387	9000	PRV57070	2	Inpatie
1	BENE11017	CLM31237	14000	PRV54750	2	Inpatie
2	BENE11026	CLM78930	2000	PRV53758	1	Inpatie
3	BENE11031	CLM56810	16000	PRV55825	2	Inpatie
4	BENE11085	CLM34625	19000	PRV52338	1	Inpatie



```

In [14]: # Now we will merge the claims table with beneficiary

# Here we will create the chronic condition col.

chronic_cols = [ "ChronicCond_Alzheimer",
                  "ChronicCond_Heartfailure",
                  "ChronicCond_KidneyDisease",
                  "ChronicCond_Cancer",
                  "ChronicCond_ObstrPulmonary",
                  "ChronicCond_Depression",
                  "ChronicCond_Diabetes",
                  "ChronicCond_IschemicHeart",
                  "ChronicCond_Osteoporosis",
                  "ChronicCond_rheumatoidarthritis",
                  "ChronicCond_stroke" ]

# This is our final summary table after pulling data from all the tables

df = claims.merge(beneficiary[["BeneID", "Gender", "Race"] + chronic_cols], on =
df.head()

```

Out[14]:

	BenelD	ClaimID	InscClaimAmtReimbursed	Provider	NumPhysicians	ClaimTy
0	BENE11014	CLM67387	9000	PRV57070	2	Inpatie
1	BENE11017	CLM31237	14000	PRV54750	2	Inpatie
2	BENE11026	CLM78930	2000	PRV53758	1	Inpatie
3	BENE11031	CLM56810	16000	PRV55825	2	Inpatie
4	BENE11085	CLM34625	19000	PRV52338	1	Inpatie

In [15]: *# Now we need to create a col that flags claims*
It wil tell us if the claims are approved or not

```
df["ClaimApproved"] = df["InscClaimAmtReimbursed"].apply(lambda x: 1 if x > 0 else 0)
```

In [16]: `df["ClaimApproved"].value_counts()`

Out[16]:

	count
ClaimApproved	
1	130469
0	4923

dtype: int64

In [17]: *# Here we can see the distribution of Approved (x = 1) and not (x = 0) claims*

In [18]: `df.head()`

Out[18]:

	BenelD	ClaimID	InscClaimAmtReimbursed	Provider	NumPhysicians	ClaimTy
0	BENE11014	CLM67387	9000	PRV57070	2	Inpatie
1	BENE11017	CLM31237	14000	PRV54750	2	Inpatie
2	BENE11026	CLM78930	2000	PRV53758	1	Inpatie
3	BENE11031	CLM56810	16000	PRV55825	2	Inpatie
4	BENE11085	CLM34625	19000	PRV52338	1	Inpatie

In [19]: *# Now we need to do some EDA on this dataset*

In [20]: `df["Gender"].value_counts()`

Out[20]:

	count
Gender	
2	78017
1	57375

dtype: int64

In [21]: *# Here in the Gender col (1 = Male and 2 = Female)*

EDA

Univariarte Analysis


```
In [22]: Male = len(df[df["Gender"] == 1])
Female = len(df[df["Gender"] == 2])

print("Percentage of Male patient is {:.2f}%".format(Male/len(df.Gender)*100))
print("Percentage of Female patient is {:.2f}%".format(Female/len(df.Gender)*100))
```

Percentage of Male patient is 42.38%

Percentage of Female patient is 57.62%

```
In [23]: sns.countplot(x = "Gender", data = df)
plt.title("Count of Male and Female Patients")
plt.show()
```



In [24]: *# In the above graph we can see there are more "Female" patients than male.*

In [25]: `df["ClaimApproved"].value_counts()`

Out[25]: **count**

ClaimApproved

1 130469

0 4923

dtype: int64

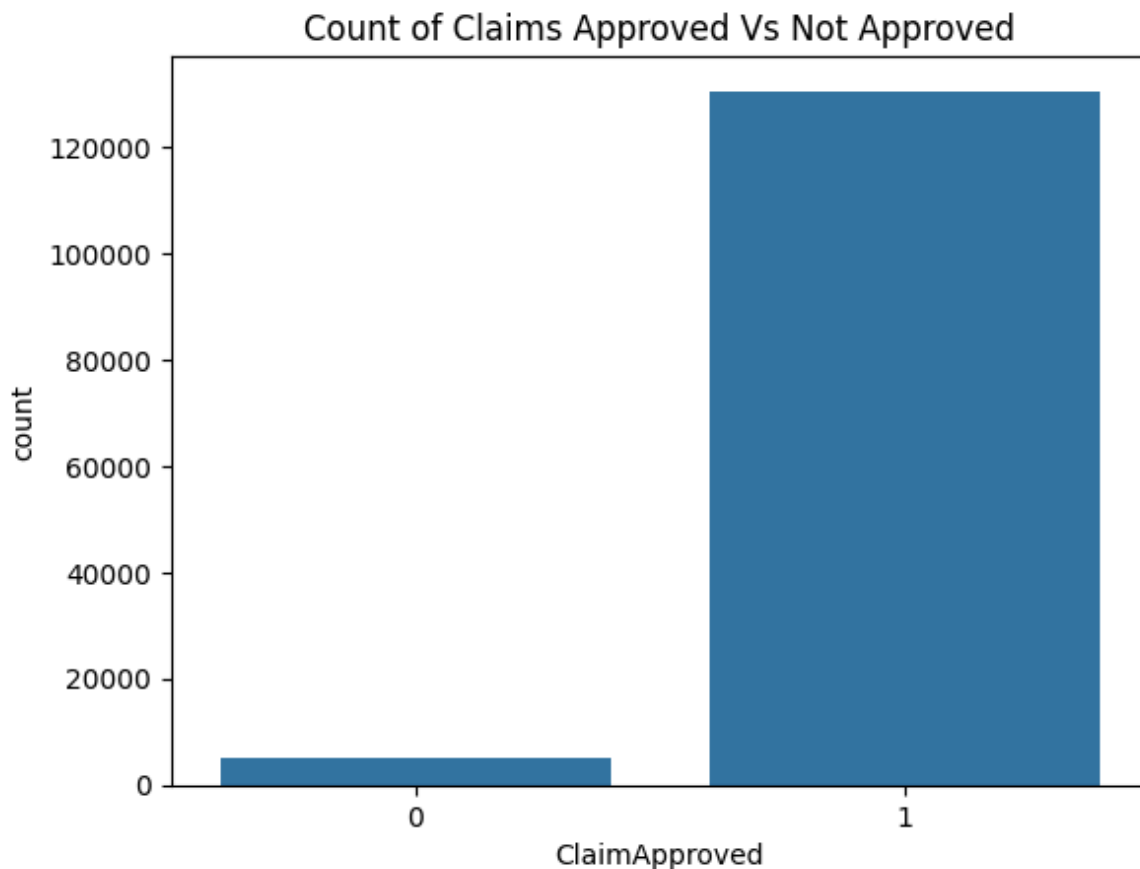
```
In [26]: claims_approved = len(df[df["ClaimApproved"] == 1])
claims_not_approved = len(df[df["ClaimApproved"] == 2])

print("Percentage of claims approved is {:.2f}%".format(Male/len(df.Gender)*100))
print("Percentage of claims not approved is {:.2f}%".format(Female/len(df.Gender)
```

Percentage of claims approved is 42.38%

Percentage of claims not approved is 57.62%

```
In [27]: sns.countplot(x = "ClaimApproved", data = df)
plt.title("Count of Claims Approved Vs Not Approved")
plt.show()
```



In [27]:

In [28]: `df["Race"].value_counts()`

Out[28]:

	count
Race	
1	115067
2	13496
3	4118
5	2711

dtype: int64In [29]: *# Percentage of different races r1, r2, r3, r4*

```

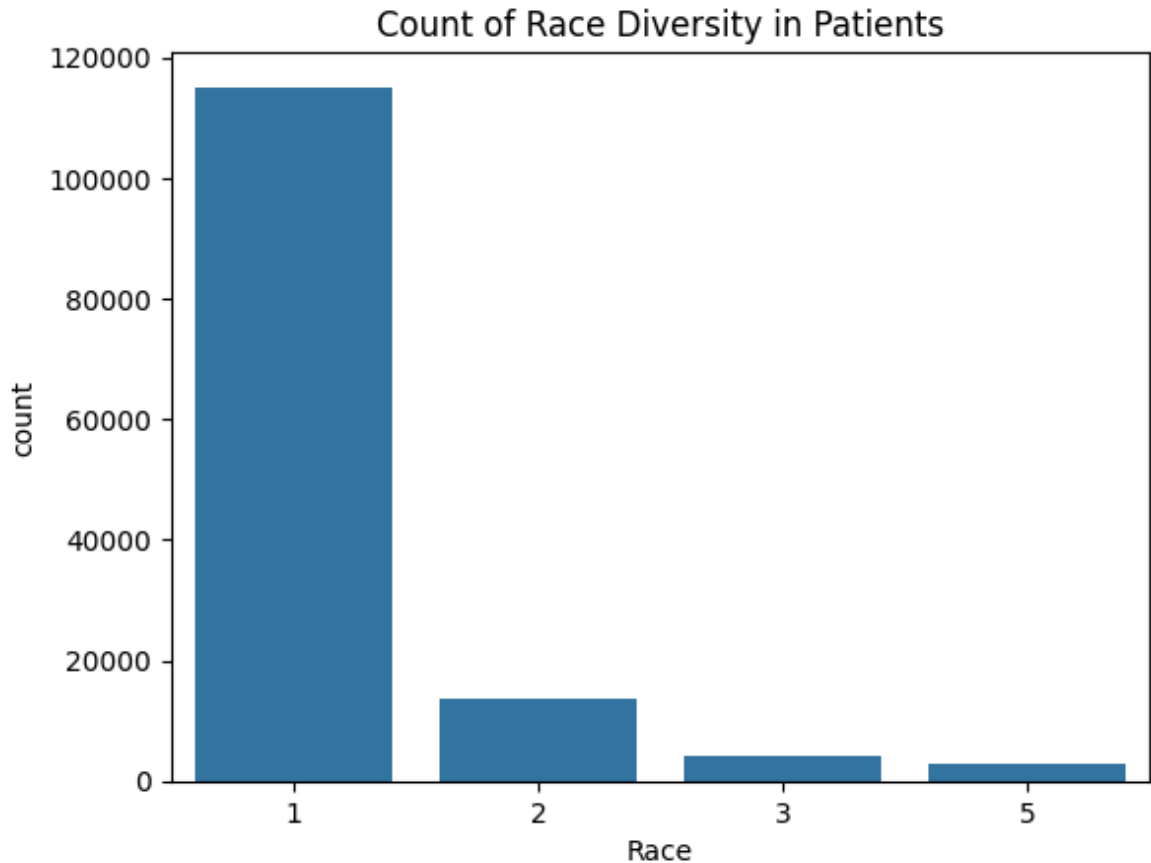
r1 = len(df[df["Race"] == 1])
r2 = len(df[df["Race"] == 2])
r3 = len(df[df["Race"] == 3])
r4 = len(df[df["Race"] == 4])

print("Percentage of r1 patient is {:.2f}%".format(r1/len(df.Race)*100))
print("Percentage of r2 patient is {:.2f}%".format(r2/len(df.Race)*100))
print("Percentage of r3 patient is {:.2f}%".format(r3/len(df.Race)*100))
print("Percentage of r4 patient is {:.2f}%".format(r4/len(df.Race)*100))

```


Percentage of r1 patient is 84.99%
 Percentage of r2 patient is 9.97%
 Percentage of r3 patient is 3.04%
 Percentage of r4 patient is 0.00%

```
In [30]: sns.countplot(x = "Race", data = df)
plt.title("Count of Race Diversity in Patients")
plt.show()
```



```
In [31]: df.describe()
```

```
Out[31]:
```

	InscClaimAmtReimbursed	NumPhysicians	Gender	Race	Chronic
count	135392.000000	135392.000000	135392.000000	135392.000000	
mean	981.307906	1.557478	1.576231	1.240605	
std	3788.177532	0.636914	0.494157	0.695578	
min	0.000000	0.000000	1.000000	1.000000	
25%	40.000000	1.000000	1.000000	1.000000	
50%	80.000000	1.000000	2.000000	1.000000	
75%	300.000000	2.000000	2.000000	1.000000	
max	125000.000000	3.000000	2.000000	5.000000	




```
In [32]: plt.figure(figsize=(15, 10))
cols = ['InscClaimAmtReimbursed']
```

```

for i in range(len(cols)):
    plt.subplot(3, 3, i + 1) # 3 rows x 3 cols grid
    sns.histplot(df[cols[i]], kde = True, bins = 30)
    plt.title(f'Distribution of {cols[i]}')

plt.tight_layout()
plt.show()

```



In [32]:

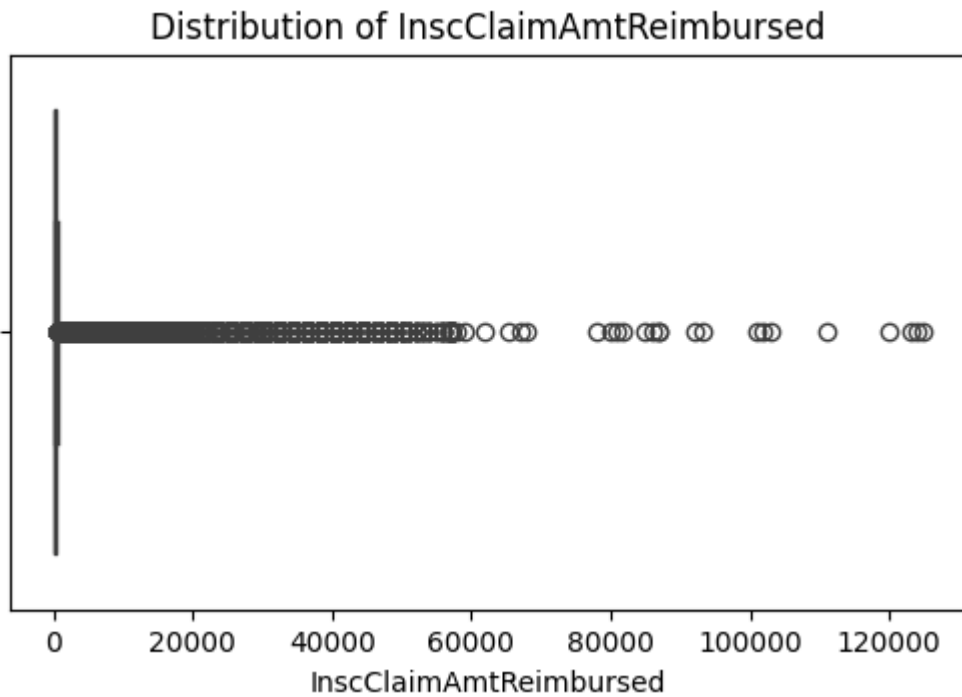
```

plt.figure(figsize=(15, 10))
cols = ['InscClaimAmtReimbursed']

for i in range(len(cols)):
    plt.subplot(3, 3, i + 1) # 3 rows x 3 cols grid
    sns.boxplot(x = df[cols[i]])
    plt.title(f'Distribution of {cols[i]}')

plt.tight_layout()
plt.show()

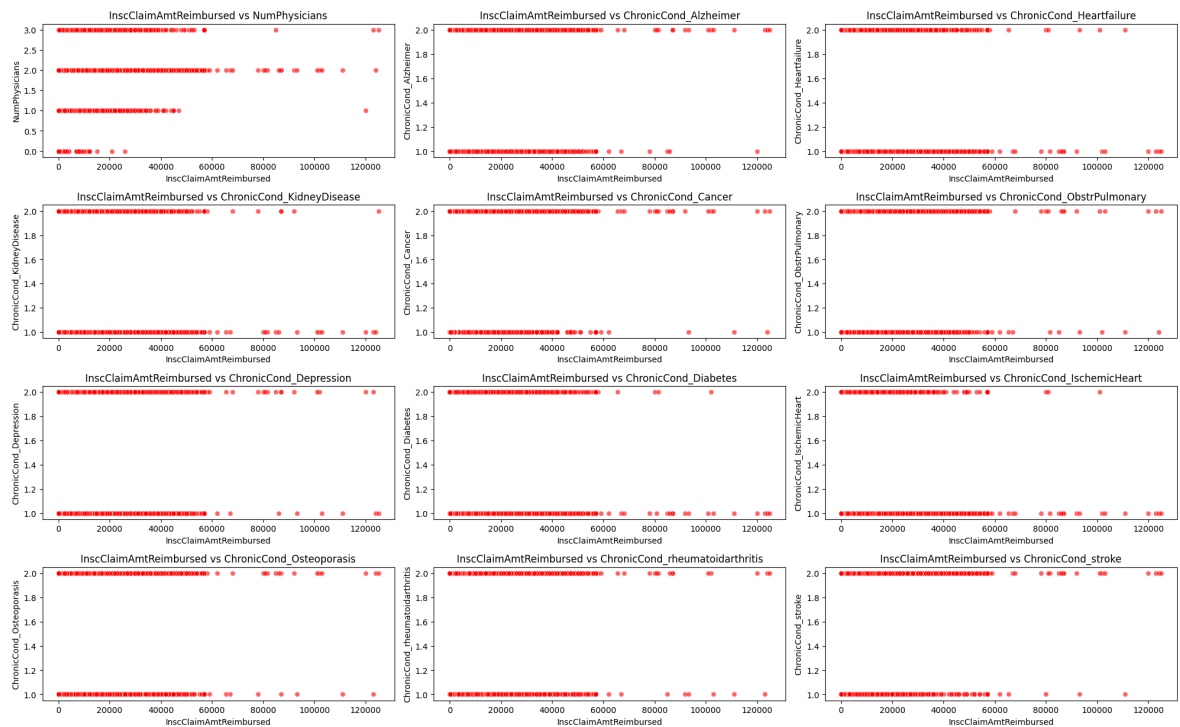
```



```
In [34]: plt.figure(figsize=(20, 15))
cols = [
    'NumPhysicians',
    'ChronicCond_Alzheimer',
    'ChronicCond_Heartfailure',
    'ChronicCond_KidneyDisease',
    'ChronicCond_Cancer',
    'ChronicCond_ObstrPulmonary',
    'ChronicCond_Depression',
    'ChronicCond_Diabetes',
    'ChronicCond_IschemicHeart',
    'ChronicCond_Osteoporosis',
    'ChronicCond_rheumatoidarthritis',
    'ChronicCond_stroke'
]

for i in range(len(cols)):
    plt.subplot(5, 3, i + 1) # 3 rows x 3 cols grid
    sns.scatterplot(data = df, x = 'InscClaimAmtReimbursed', y = cols[i], color
    plt.title(f'InscClaimAmtReimbursed vs {cols[i]}')

plt.tight_layout()
plt.show()
```



In [34]:


Multivariate Analysis

In [35]:

```
# Correlation Table - Heat Map
plt.figure(figsize=(20, 15))

cols = ['InscClaimAmtReimbursed',
        'ClaimApproved',
        'NumPhysicians',
        'ChronicCond_Alzheimer',
        'ChronicCond_Heartfailure',
        'ChronicCond_KidneyDisease',
        'ChronicCond_Cancer',
        'ChronicCond_ObstrPulmonary',
        'ChronicCond_Depression',
        'ChronicCond_Diabetes',
        'ChronicCond_IschemicHeart',
        'ChronicCond_Osteoporosis',
        'ChronicCond_rheumatoidarthritis',
        'ChronicCond_stroke']

sns.heatmap(data = df[cols].corr(), annot = True, cmap = 'rocket')
plt.show()
```



In [35]:

Data Preprocessing

In [36]: `model_data = df.copy()`In [37]: `# Removing unneseccary cols`
`model_data.columns`Out[37]: `Index(['BeneID', 'ClaimID', 'InscClaimAmtReimbursed', 'Provider',
 'NumPhysicians', 'ClaimType', 'Gender', 'Race', 'ChronicCond_Alzheimer',
 'ChronicCond_Heartfailure', 'ChronicCond_KidneyDisease',
 'ChronicCond_Cancer', 'ChronicCond_ObstrPulmonary',
 'ChronicCond_Depression', 'ChronicCond_Diabetes',
 'ChronicCond_IschemicHeart', 'ChronicCond_Osteoporosis',
 'ChronicCond_rheumatoidarthritis', 'ChronicCond_stroke',
 'ClaimApproved'],
 dtype='object')`In [38]: `model_data = model_data.drop(["BeneID", "ClaimID"], axis=1)`In [39]: `model_data.columns`

```
Out[39]: Index(['InscClaimAmtReimbursed', 'Provider', 'NumPhysicians', 'ClaimType',
              'Gender', 'Race', 'ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
              'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
              'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression',
              'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
              'ChronicCond_Osteoporasis', 'ChronicCond_rheumatoidarthritis',
              'ChronicCond_stroke', 'ClaimApproved'],
              dtype='object')
```

```
In [40]: # Removing Null Values
         model_data.isnull().sum()
```

```
Out[40]:
```

InscClaimAmtReimbursed	0
Provider	0
NumPhysicians	0
ClaimType	0
Gender	0
Race	0
ChronicCond_Alzheimer	0
ChronicCond_Heartfailure	0
ChronicCond_KidneyDisease	0
ChronicCond_Cancer	0
ChronicCond_ObstrPulmonary	0
ChronicCond_Depression	0
ChronicCond_Diabetes	0
ChronicCond_IschemicHeart	0
ChronicCond_Osteoporasis	0
ChronicCond_rheumatoidarthritis	0
ChronicCond_stroke	0
ClaimApproved	0

dtype: int64

```
In [41]: model_data.columns
```

```
Out[41]: Index(['InscClaimAmtReimbursed', 'Provider', 'NumPhysicians', 'ClaimType',
              'Gender', 'Race', 'ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
              'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
              'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression',
              'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
              'ChronicCond_Osteoporasis', 'ChronicCond_rheumatoidarthritis',
              'ChronicCond_stroke', 'ClaimApproved'],
              dtype='object')
```

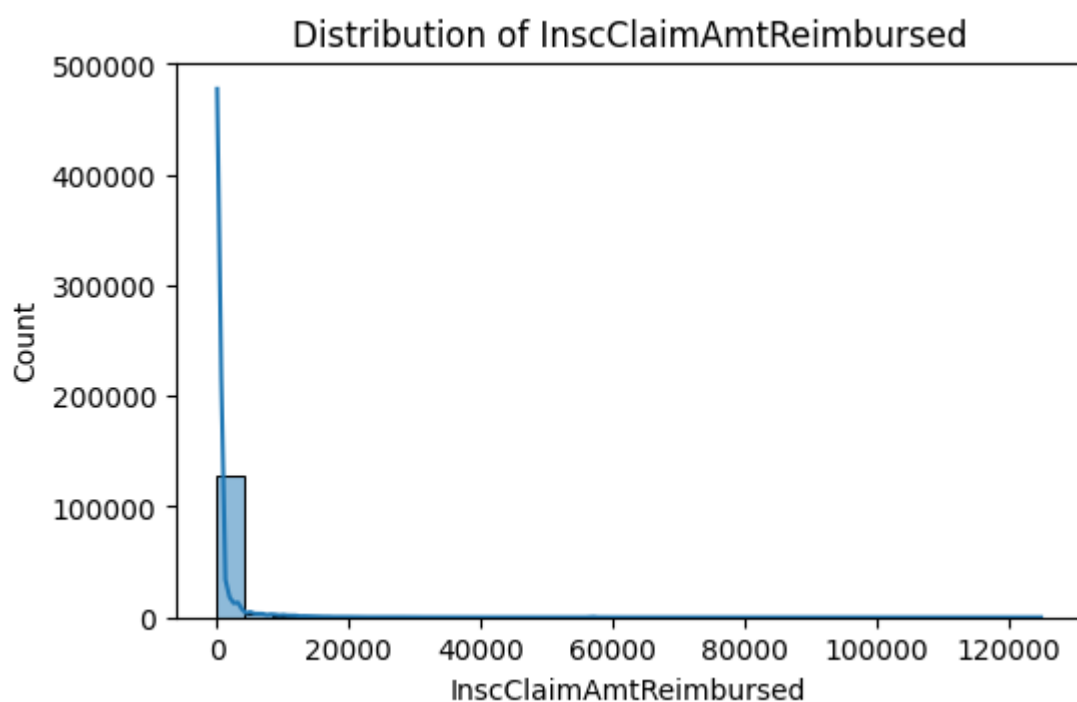
Data Transformation

```
In [42]: model_data["InscClaimAmtReimbursed"] = np.log1p(model_data["InscClaimAmtReimburs
```

```
In [43]: plt.figure(figsize=(15, 10))
cols = ['InscClaimAmtReimbursed']

for i in range(len(cols)):
    plt.subplot(3, 3, i + 1) # 3 rows x 3 cols grid
    sns.histplot(df[cols[i]], kde = True, bins = 30)
    plt.title(f'Distribution of {cols[i]}')


plt.tight_layout()
plt.show()
```



```
In [44]: plt.figure(figsize=(15, 10))
cols = ['InscClaimAmtReimbursed']

for i in range(len(cols)):
    plt.subplot(3, 3, i + 1) # 3 rows x 3 cols grid
    sns.histplot(model_data[cols[i]], kde = True, bins = 30)
    plt.title(f'Distribution of {cols[i]}')

plt.tight_layout()
plt.show()
```



In [45]: *# We can see that earlier the data did not had a normal distribution now it foll*


In [48]: `model_data = pd.get_dummies(model_data, dtype = int)`

In [49]: *# Feature separation - x and y (Target variable and independent variable)*
`x = model_data.drop(["ClaimApproved"], axis =1)`
`y = model_data["ClaimApproved"]`

In [51]: `x,y = SMOTE().fit_resample(x,y)`

In [52]: `sns.countplot(y = y, data = model_data)`

Out[52]: `<Axes: xlabel='count', ylabel='ClaimApproved'>`



```
In [50]: x = MinMaxScaler().fit_transform(x)
x
```

```
Out[50]: array([[0.77582066, 0.66666667, 1.          , ..., 0.          , 1.          ,
                  0.          ],
                [0.81346467, 0.66666667, 1.          , ..., 0.          , 1.          ,
                  0.          ],
                [0.64769533, 0.33333333, 0.          , ..., 0.          , 1.          ,
                  0.          ],
                ...,
                [0.29260094, 0.33333333, 1.          , ..., 0.          , 0.          ,
                  1.          ],
                [0.29260094, 0.66666667, 0.          , ..., 0.          , 0.          ,
                  1.          ],
                [0.66670029, 0.33333333, 0.          , ..., 0.          , 0.          ,
                  1.          ]])
```

```
In [53]: #Train test split
```

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.2, random_sta
```

```
In [54]: #Modelling:
```

```
#Logistic regression
```

```
classifier = LogisticRegression(solver = 'saga', max_iter = 500, random_state= 1
```

```
classifier.fit(x_train, y_train)
```

```
y_pred = classifier.predict(x_test)
```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	26114
1	1.00	1.00	1.00	26074
accuracy			1.00	52188
macro avg	1.00	1.00	1.00	52188
weighted avg	1.00	1.00	1.00	52188

```
In [55]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report

# Create the model
dt_classifier = DecisionTreeClassifier(random_state=1)

# Train
dt_classifier.fit(x_train, y_train)

# Predict
y_pred_dt = dt_classifier.predict(x_test)

# Evaluation
print("Decision Tree Classification Report:")
print(classification_report(y_test, y_pred_dt))
```

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	26114
1	1.00	1.00	1.00	26074
accuracy			1.00	52188
macro avg	1.00	1.00	1.00	52188
weighted avg	1.00	1.00	1.00	52188

```
In [56]: from sklearn.ensemble import RandomForestClassifier

# Create the model
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=1)

# Train
rf_classifier.fit(x_train, y_train)

# Predict
y_pred_rf = rf_classifier.predict(x_test)

# Evaluation
print("Random Forest Classification Report:")
print(classification_report(y_test, y_pred_rf))
```

Random Forest Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	26114
1	1.00	1.00	1.00	26074
accuracy			1.00	52188
macro avg	1.00	1.00	1.00	52188
weighted avg	1.00	1.00	1.00	52188

In []: